

```
In [1]: import pandas as pd
import seaborn as sns
```

```
In [2]: data = pd.read_csv('diabetes.csv')
```

```
In [3]: data.head()
```

Out[3]:	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
4   Insulin                768 non-null    int64
5   BMI                   768 non-null    float64
6   Pedigree               768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [5]: data.shape
```

Out[5]: (768, 9)

```
In [6]: #Check for null or missing values
data.isnull().sum()
```


```
Out[6]: Pregnancies      0
         Glucose          0
         BloodPressure    0
         SkinThickness    0
         Insulin          0
         BMI              0
         Pedigree         0
         Age              0
         Outcome          0
         dtype: int64
```

```
In [7]: #input data
x= data.drop('Outcome',axis=1)

#output data
y=data['Outcome']
```

```
In [8]: sns.countplot(x=y)
```

```
Out[8]: <AxesSubplot:xlabel='Outcome', ylabel='count'>
```



```
In [9]: y.value_counts()
```

```
Out[9]: 0    500
        1    268
        Name: Outcome, dtype: int64
```

```
In [10]: #Feature Scaling---->feature are scaled in one range (0-1)
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()
x_scaled=scaler.fit_transform(x)
```

```
In [11]: x_scaled
```

```
Out[11]: array([[0.35294118, 0.74371859, 0.59016393, ..., 0.50074516, 0.23441503,
                  0.48333333],
                 [0.05882353, 0.42713568, 0.54098361, ..., 0.39642325, 0.11656704,
                  0.16666667],
                 [0.47058824, 0.91959799, 0.52459016, ..., 0.34724292, 0.25362938,
                  0.18333333],
                 ...,
                 [0.29411765, 0.6080402 , 0.59016393, ..., 0.390462 , 0.07130658,
                  0.15          ],
                 [0.05882353, 0.63316583, 0.49180328, ..., 0.4485842 , 0.11571307,
                  0.43333333],
                 [0.05882353, 0.46733668, 0.57377049, ..., 0.45305514, 0.10119556,
                  0.03333333]])
```

```
In [12]: #Cross Validation  
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.25, random_
```

```
In [13]: x.shape
```

```
Out[13]: (768, 8)
```

```
In [14]: x_train.shape
```

```
Out[14]: (576, 8)
```

```
In [15]: x_test.shape
```

```
Out[15]: (192, 8)
```

K-Nearest Neighbors algorithm

```
In [16]: #KNN  
#import class  
from sklearn.neighbors import KNeighborsClassifier
```

```
In [17]: #Create the object---num of neighbours 2  
knn=KNeighborsClassifier(n_neighbors= 33)
```

```
In [18]: #Train the algorithm  
knn.fit(x_train,y_train)
```


```
Out[18]: KNeighborsClassifier(n_neighbors=33)
```

```
In [19]: #Predict on test data  
y_pred=knn.predict(x_test)
```

```
In [20]: #import the evaluation matrix  
from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score, classification_repor
```

```
In [21]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[21]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2037377d3d0>
```



```
In [22]: accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of Logistic Regression: {accuracy}")
```

Accuracy of Logistic Regression: 0.7916666666666666

```
In [23]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.79	0.94	0.86	130
1	0.79	0.48	0.60	62
accuracy			0.79	192
macro avg	0.79	0.71	0.73	192
weighted avg	0.79	0.79	0.78	192

Naive Bayes algorithm

```
In [24]: from sklearn.naive_bayes import GaussianNB
```

```
In [25]: # Create a Naive Bayes classifier
naive_bayes = GaussianNB()
```


```
In [26]: # Train the algorithm
naive_bayes.fit(x_train, y_train)
```

```
Out[26]: GaussianNB()
```

```
In [27]: # Predict on test data
y_pred = naive_bayes.predict(x_test)
```

```
In [28]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[28]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x203733ec820>
```



```
In [29]: # Evaluate the performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of Naive Bayes: {accuracy}")
```

Accuracy of Naive Bayes: 0.765625

```
In [30]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.88	0.84	130
1	0.67	0.53	0.59	62
accuracy			0.77	192
macro avg	0.74	0.70	0.71	192
weighted avg	0.76	0.77	0.76	192

Decision Tree algorithm

```
In [31]: from sklearn.tree import DecisionTreeClassifier
```

```
In [32]: # Create a Decision Tree classifier
decision_tree = DecisionTreeClassifier()
```


```
In [33]: # Train the algorithm
decision_tree.fit(x_train, y_train)
```

```
Out[33]: DecisionTreeClassifier()
```

```
In [34]: # Predict on test data
y_pred = decision_tree.predict(x_test)
```

```
In [35]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[35]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20373b57580>
```



```
In [36]: # Evaluate the performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of Decision Tree: {accuracy}")
```

Accuracy of Decision Tree: 0.7447916666666666

```
In [37]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.82	0.81	130
1	0.61	0.60	0.60	62
accuracy			0.74	192
macro avg	0.71	0.71	0.71	192
weighted avg	0.74	0.74	0.74	192

Random Forest algorithm

```
In [38]: from sklearn.ensemble import RandomForestClassifier
```

```
In [39]: # Create a Random Forest classifier
random_forest = RandomForestClassifier(n_estimators=100) # You can adjust the number of
```


```
In [40]: # Train the algorithm
random_forest.fit(x_train, y_train)
```

```
Out[40]: RandomForestClassifier()
```

```
In [41]: # Predict on test data
y_pred = random_forest.predict(x_test)
```

```
In [42]: ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[42]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x20373a6db50>
```



```
In [43]: # Evaluate the performance
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of Random Forest: {accuracy}")
```

Accuracy of Random Forest: 0.7708333333333334

```
In [44]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.88	0.84	130
1	0.68	0.55	0.61	62
accuracy			0.77	192
macro avg	0.74	0.71	0.72	192
weighted avg	0.76	0.77	0.76	192